

Group No: 53

Group Member Names:

Name	BITS ID	Contribution
Mujtaba Rasool	2023ac05819	100%
Kartik Batta	2023ac05422	100%
Sachin Laddha	2023ac05564	100%

1.Problem statement:

- Develop a reinforcement learning agent using dynamic programming to solve the Treasure Hunt problem in a FrozenLake environment. The agent must learn the optimal policy for navigating the lake while avoiding holes and maximizing its treasure collection.

2.Scenario:

- A treasure hunter is navigating a slippery 5x5 FrozenLake grid. The objective is to navigate through the lake collecting treasures while avoiding holes and ultimately reaching the exit (goal). Grid positions on a 5x5 map with tiles labeled as S, F, H, G, T. The state includes the current position of the agent and whether treasures have been collected.

Objective

- The agent must learn the optimal policy π^* using dynamic programming to maximize its cumulative reward while navigating the lake.

About the environment

The environment consists of several types of tiles:

- Start (S): The initial position of the agent, safe to step.
- Frozen Tiles (F): Frozen surface, safe to step.
- Hole (H): Falling into a hole ends the game immediately (die, end).
- Goal (G): Exit point; reaching here ends the game successfully (safe, end).
- Treasure Tiles (T): Added to the environment. Stepping on these tiles awards +5 reward but does not end the game.

After stepping on a treasure tile, it becomes a frozen tile (F). The agent earns rewards as follows:

- Reaching the goal (G): +10 reward.
- Falling into a hole (H): -10 reward.
- Collecting a treasure (T): +5 reward.
- Stepping on a frozen tile (F): 0 reward.

States

- Current position of the agent (row, column).
- A boolean flag (or equivalent) for whether each treasure has been collected.

Actions

- Four possible moves: up, down, left, right

Rewards

- Goal (G): +10.
- Treasure (T): +5 per treasure.
- Hole (H): -10.
- Frozen tiles (F): 0.

Environment

Modify the FrozenLake environment in OpenAI Gym to include treasures (T) at certain positions. Inherit the original FrozenLakeEnv and modify the reset and step methods accordingly. Example grid:

S	F	F	H	T
F	H	F	F	F
F	F	F	T	F
T	F	H	F	F
F	F	F	F	G

Expected Outcomes:

1. Create the custom environment by modifying the existing "FrozenLakeNotSlippery-v0" in OpenAI Gym and Implement the dynamic programming using value iteration and policy improvement to learn the optimal policy for the Treasure Hunt problem.
2. Calculate the state-value function (V^*) for each state on the map after learning the optimal policy.
3. Compare the agent's performance with and without treasures, discussing the trade-offs in reward maximization.

4. Visualize the agent's direction on the map using the learned policy.
5. Calculate expected total reward over multiple episodes to evaluate performance.

Import required libraries and Define the custom environment - 2 Marks

```
In [1]: pip install gym numpy
```

```
Requirement already satisfied: gym in /Users/sachinladdha/anaconda3/lib/python3.12/site-packages (0.26.2)
Requirement already satisfied: numpy in /Users/sachinladdha/anaconda3/lib/python3.12/site-packages (1.26.4)
Requirement already satisfied: cloudpickle>=1.2.0 in /Users/sachinladdha/anaconda3/lib/python3.12/site-packages (from gym) (3.0.0)
Requirement already satisfied: gym-notices>=0.0.4 in /Users/sachinladdha/anaconda3/lib/python3.12/site-packages (from gym) (0.0.8)
Note: you may need to restart the kernel to use updated packages.
```

```
In [2]: # Import statements
import gym
from gym.envs.toy_text import FrozenLakeEnv
import numpy as np
```

```
In [3]: # Custom environment to create the given grid and respective functions that are req

#Include functions to take an action, get reward, to check if episode is over
class FrozenLakeTreasureEnv(FrozenLakeEnv):
    def __init__(self):
        super().__init__(desc=np.asarray([
            [b'S', b'F', b'F', b'H', b'T'],
            [b'F', b'H', b'F', b'F', b'F'],
            [b'F', b'F', b'F', b'T', b'F'],
            [b'T', b'F', b'H', b'F', b'F'],
            [b'F', b'F', b'F', b'F', b'G']
        ], dtype='|S1'), is_slippery=False)

        self.treasure_locations = [(0, 4), (2, 3), (3, 0)]
        self.treasures_collected = set()

        self.nS = self.observation_space.n
        self.nA = self.action_space.n
        self.ncol = 5

    def reset(self):
        self.s = 0
        self.treasures_collected = set()
        return self.s

    def step(self, a):
        self.lastaction = a
        curr_pos = (self.s // self.ncol, self.s % self.ncol)
        next_state = self._move(self.s, a)
        next_pos = (next_state // self.ncol, next_state % self.ncol)
```

```

reward = 0
done = False
next_tile = self.desc[next_pos[0]][next_pos[1]]
if next_tile == b'G':
    reward = 10
    done = True
elif next_tile == b'H':
    reward = -10
    done = True
elif next_tile == b'F':
    reward = 0

if next_pos in self.treasure_locations and next_pos not in self.treasures_c

    reward += 5
    self.treasures_collected.add(next_pos)
    self._update_desc()

self.s = next_state
return next_state, reward, done, {}

def _move(self, state, action):
    row = state // self.ncol
    col = state % self.ncol

    if action == 0: # left
        col = max(col - 1, 0)
    elif action == 1: # down
        row = min(row + 1, self.ncol - 1)
    elif action == 2: # right
        col = min(col + 1, self.ncol - 1)
    elif action == 3: # up
        row = max(row - 1, 0)

    return row * self.ncol + col

def _update_desc(self):
    desc = self.desc.tolist()
    for treasure in self.treasures_collected:
        row, col = treasure
        desc[row][col] = b'F'
    self.desc = np.asarray(desc, dtype='|S1')

```

In []:

Value Iteration Algorithm - 1 Mark

```

In [4]: def value_iteration(env, gamma=0.9, theta=1e-8):
V = np.zeros(env.observation_space.n)
policy = np.zeros(env.observation_space.n, dtype=int)

while True:
    delta = 0
    for s in range(env.observation_space.n):
        v = V[s]

```

```

        action_values = np.zeros(env.action_space.n)
        actions = env.P[s] # Use environment's transition probabilities
        for a in range(env.action_space.n):
            action_value = 0
            for prob, next_state, reward, done in actions[a]:
                row, col = next_state // env.ncol, next_state % env.ncol
                tile = env.desc[row][col]

                if tile == b'G':
                    reward = 10
                elif tile == b'H':
                    reward = -10
                elif tile == b'T':
                    reward = 5
                else: # Frozen tile
                    reward = 0

                action_value += prob * (reward + gamma * V[next_state])
            action_values[a] = action_value

        best_action = np.argmax(action_values)
        V[s] = action_values[best_action]
        policy[s] = best_action
        delta = max(delta, np.abs(v - V[s]))
    if delta < theta:
        break
    return V, policy

```

Policy Improvement Function - 1 Mark

In []:

Print the Optimal Value Function

```

In [5]: def print_value_function(V, env):
        print("\nOptimal Value Function:")
        map_size = int(np.sqrt(env.observation_space.n))
        for i in range(map_size):
            for j in range(map_size):
                print("{:6.2f}".format(V[i * map_size + j]), end=" ")
            print()

```

Visualization of the learned optimal policy - 1 Mark

```

In [6]: def visualize_policy(policy, map_size):
        actions = ["←", "↓", "→", "↑"]
        policy_grid = policy.reshape(map_size, map_size)
        print("\nOptimal Policy:")
        for row in policy_grid:
            print(" ".join([actions[action] for action in row]))

```

Evaluate the policy - 1 Mark

```
In [7]: def evaluate_policy(env, policy, num_episodes=100):
        total_rewards = []
        for _ in range(num_episodes):
            state = env.reset()
            if isinstance(state, tuple):
                state = state[0]

            total_reward = 0
            done = False
            while not done:
                action = policy[int(state)]
                step_result = env.step(action)

                if len(step_result) == 5:
                    next_state, reward, terminated, truncated, info = step_result
                    done = terminated or truncated
                else:
                    next_state, reward, done, info = step_result

                if isinstance(next_state, tuple):
                    next_state = next_state[0]

                total_reward += reward
                state = next_state
            total_rewards.append(total_reward)
        return np.mean(total_rewards)
```

Main Execution

```
In [8]: # Create environment with treasures
        env = FrozenLakeTreasureEnv()

        # Run value iteration
        V, policy = value_iteration(env)

        # Print results for environment with treasures
        print("\n=== Environment With Treasures ===")
        print_value_function(V, env)
        visualize_policy(policy, 5)
        reward_with_treasures = evaluate_policy(env, policy)
        print("\nAverage Reward with Treasures:", reward_with_treasures)

        # Create environment without treasures
        env_no_treasures = FrozenLakeEnv(desc=np.asarray([
            [b'S', b'F', b'F', b'H', b'F'],
            [b'F', b'H', b'F', b'F', b'F'],
            [b'F', b'F', b'F', b'F', b'F'],
            [b'F', b'F', b'H', b'F', b'F'],
            [b'F', b'F', b'F', b'F', b'G']
        ], dtype='|S1'), is_slippery=False)
```

```

# Run value iteration for environment without treasures
V_no_treasures, policy_no_treasures = value_iteration(env_no_treasures)

# Print results for environment without treasures
print("\n=== Environment Without Treasures ===")
print_value_function(V_no_treasures, env_no_treasures)
visualize_policy(policy_no_treasures, 5)
reward_no_treasures = evaluate_policy(env_no_treasures, policy_no_treasures)
print("\nAverage Reward without Treasures:", reward_no_treasures)

```

=== Environment With Treasures ===

Optimal Value Function:

```

51.88  56.79  63.10 -100.00  72.90
57.64 -100.00  70.11  77.90  81.00
64.05  70.11  77.90  81.00  90.00
65.61  72.90 -100.00  90.00 100.00
72.90  81.00  90.00 100.00 100.00

```

Optimal Policy:

```

↓ → ↓ ← ↓
↓ ← ↓ ↓ ↓
↓ → → ↓ ↓
↓ ↓ ← ↓ ↓
→ → → → ←

```

Average Reward with Treasures: 15.0

=== Environment Without Treasures ===

Optimal Value Function:

```

47.83  53.14  59.05 -100.00  72.90
53.14 -100.00  65.61  72.90  81.00
59.05  65.61  72.90  81.00  90.00
65.61  72.90 -100.00  90.00 100.00
72.90  81.00  90.00 100.00 100.00

```

Optimal Policy:

```

↓ → ↓ ← ↓
↓ ← ↓ ↓ ↓
↓ ↓ → ↓ ↓
↓ ↓ ← ↓ ↓
→ → → → ←

```

Average Reward without Treasures: 1.0